

DATA MINING

/CISC 873 - Steven Ding

/Metric/Loss

Performance Metric

- Not all errors are equally important.
- It depends on what you NEED. (Need-driven)
- Classification:
 - TP, FP, TN, FN, Accuracy, Recall, Precision, Sensitivity, Specificity, AUROC, F1, ...
- Prediction (Regression)
 - MSE, MAE, MSLE, R2, ...

Confusion Matrix

Real class\Predicted class	C ₁	C ₂
C ₁	True positive (TP)	False negative (FN)
C ₂	False positive (FP)	True negative (TN)

Fact\Answer	Yes	no	total
yes	6 (TP)	3 (FN)	9 (P)
no	2 (FP)	3 (TN)	5 (N)
total	8	6	14

- Accuracy: percentage of test set tuples that are correctly classified:
 - $Accuracy = \frac{TP+TN}{P+N} = \frac{6+3}{14}$
 - $Recall = \frac{TP}{P} = \frac{6}{9}$ (percentage of positives correctly answered)
 - $Precision = \frac{TP}{TP+FP} = \frac{6}{6+2}$ (chance that a positive answer is correct)

Confusion Matrix – F1

Real class\Predicted class	C ₁	C ₂
C ₁	True positive (TP)	False negative (FN)
C ₂	False positive (FP)	True negative (TN)

Fact\Answer	Yes	no	total
yes	6 (TP)	3 (FN)	9 (P)
no	2 (FP)	3 (TN)	5 (N)
total	8	6	14

- F measure: the harmonic mean of precision and recall:


- $$F = \frac{2 * Precision * Recall}{Precision + Recall}$$

Multi-class Measures

- Micro-average
 - sum up the individual true positives, false positives, and false negatives for different classes and then apply them to get the statistics.
- Macro-average
 - take the average of the chosen metric (e.g. precision, recall etc.) of the system on different classes

AUROC

- Area under receiver operating characteristic curve



<https://stats.stackexchange.com/questions/225210/accuracy-vs-area-under-the-roc-curve>

Typical Training Loss

- Classification loss
 - Logistic/Binary Cross Entropy (Binary)
 - Hinge Loss (Binary)
 - KL divergence (Binary)
 - Cross Entropy (Categorical, one-hot encoded)
 - Sparse Categorical Cross entropy (Categorical)

Binary Cross Entropy (Binary)

$$-(y \log(p) + (1 - y) \log(1 - p))$$

```
# Compute cross entropy from probabilities.  
bce = target * math_ops.log(output + epsilon())  
bce += (1 - target) * math_ops.log(1 - output + epsilon())  
return -bce
```

```
y_true = [[0., 1.], [0., 0.]]  
y_pred = [[0.6, 0.4], [0.4, 0.6]]
```

```
bce = tf.keras.losses.BinaryCrossentropy()  
bce(y_true, y_pred).numpy()
```


Hinge Loss

$$\text{loss}(\mathbf{w}) = \sum_i^n \max(0, 1 - y_i \mathbf{w}^T \mathbf{x})$$

```
1381 y_pred = ops.convert_to_tensor_v2(y_pred)
1382 y_true = math_ops.cast(y_true, y_pred.dtype)
1383 y_true = _maybe_convert_labels(y_true)
1384 return K.mean(math_ops.maximum(1. - y_true * y_pred, 0.), axis=-1)
1385
```

```
y_true = [[0., 1.], [0., 0.]]
y_pred = [[0.6, 0.4], [0.4, 0.6]]
```

```
h = tf.keras.losses.Hinge()
h(y_true, y_pred).numpy()
```

KL divergence

- Kullback-Leibler divergence loss
- Measure how one probability distribution differs from a second
- Asymmetric

$$\text{loss} = y * \log(y / p)$$

```
.....  
y_pred = ops.convert_to_tensor_v2(y_pred)  
y_true = math_ops.cast(y_true, y_pred.dtype)  
y_true = K.clip(y_true, K.epsilon(), 1)  
y_pred = K.clip(y_pred, K.epsilon(), 1)  
return math_ops.reduce_sum(y_true * math_ops.log(y_true / y_pred), axis=-1)
```

Categorical Target

- Softmax Function:
 - logits: output from the dense layer before prediction

```
pred = tf.exp(logits) / tf.reduce_sum(tf.exp(logits))
```


```
logits = [1, 2, 3, 4, 5, 6]
```

```
pred = [  
0.00426977826282382,  
0.011606461368501186,  
0.03154963254928589,  
0.08576078712940216,  
0.23312200605869293,  
0.6336913108825684]
```

Cross Entropy (Categorical)

pred = [y = [
0.00426977826282382,	0,
0.011606461368501186,	0,
0.03154963254928589,	0,
0.08576078712940216,	0,
0.23312200605869293,	1,
0.6336913108825684]	0]

$$L = -\frac{1}{m} \sum_{i=1}^m y_i \cdot \log(\hat{y}_i)$$



Cross Entropy (TF)

 View aliases

```
tf.keras.losses.categorical_crossentropy(  
    y_true, y_pred, from_logits=False, label_smoothing=0  
)
```



Standalone usage:

```
>>> y_true = [[0, 1, 0], [0, 0, 1]]  
>>> y_pred = [[0.05, 0.95, 0], [0.1, 0.8, 0.1]]  
>>> loss = tf.keras.losses.categorical_crossentropy(y_true, y_pred)  
>>> assert loss.shape == (2,)  
>>> loss.numpy()  
array([0.0513, 2.303], dtype=float32)
```




Cross Entropy (TF)

```
tf.keras.losses.SparseCategoricalCrossentropy(  
    from_logits=False, reduction=losses_utils.ReductionV2.AUTO,  
    name='sparse_categorical_crossentropy'  
)
```

```
>>> y_true = [1, 2]  
>>> y_pred = [[0.05, 0.95, 0], [0.1, 0.8, 0.1]]  
>>> # Using 'auto'/'sum_over_batch_size' reduction type.  
>>> scce = tf.keras.losses.SparseCategoricalCrossentropy()  
>>> scce(y_true, y_pred).numpy()  
1.177
```

Summary



https://scikit-learn.org/stable/auto_examples/linear_model/plot_sgd_loss_functions.html